ASIAN INSTITUTE OF TECHNOLOGY
FACULTY OF ADVANCED SCIENCE AND TECHNOLOGY
COMPUTER SCIENCE AND COMMUNICATION TECHNOLOGY

MindTrace: Text Mining and NLP-Driven Emotion Prediction Using Machine Learning and Deep Learning Approaches

Submitted by

Aye Khin Khin Hpone (Yolanda Lim)
125970

COMPUTER SCIENCE

Asian Institute of Technology
Pathum Thani, Thailand
March 2026

ABSTRACT

Text-based emotion detection from social media has grown in importance as platforms such as Twitter generate large volumes of emotionally expressive informal text. This study proposes MindTrace, a comparative evaluation of four classification models — SVM, XGBoost, CNN, and BiLSTM — applied to a 416,809-sample Twitter corpus annotated across six emotion categories: anger, joy, sadness, fear, surprise, and love. Existing comparative studies vary widely in dataset selection, preprocessing methods, and evaluation metrics, making direct performance comparisons unreliable. This study addresses that gap by fixing all experimental conditions across all four model families on this same corpus. A standardised NLP preprocessing pipeline handles tokenisation, stopword removal with negation preservation, lemmatisation, and class balancing. ML models use TF-IDF vectorisation; DL models use trainable embedding layers. Performance is assessed using macro-F1, weighted-F1, and class-wise confusion matrices to account for the significant class imbalance (imbalance ratio ≈ 9.5:1). The findings are expected to provide empirically grounded guidance for architecture selection in emotion classification tasks, with applications in social media analytics, mental health monitoring, and emotion-aware AI systems.

Keywords: *emotion classification, natural language processing, machine learning, deep learning, BiLSTM, CNN, SVM, XGBoost, TF-IDF, Twitter, multi-class classification, social media analytics*

1. INTRODUCTION

1.1 Background and Motivation

Emotion detection from textual data has become increasingly significant in the digital era, where individuals frequently express their psychological states and reactions through social media platforms. Automated emotion recognition systems enable deeper understanding of user behaviour, going beyond traditional sentiment polarity analysis. Such systems have practical applications in marketing analytics, customer relationship management, healthcare monitoring, and human-computer interaction (Nandwani & Verma, 2021; Poria et al., 2019).Unlike sentiment analysis, which categorises content as positive, negative, or neutral, emotion detection aims to identify discrete emotional states such as anger, joy, sadness, fear, surprise, and love. However, accurately classifying emotions remains challenging due to contextual ambiguity, informal language usage, sarcasm, and class imbalance within datasets (Seyeditabari et al., 2018). This proposed study intends to develop a structured NLP pipeline and conduct a comparative evaluation between traditional machine learning and deep learning approaches for multi-class emotion classification.

1.2 Problem Statement

Despite significant research progress, a fundamental methodological gap persists. Comparative analyses across studies frequently differ in dataset selection, preprocessing strategies, feature engineering approaches, and evaluation metrics — making observed performance differences attributable to experimental inconsistencies rather than intrinsic model capabilities. Multi-class emotion classification across semantically similar categories (e.g., fear vs. surprise, sadness vs. love) remains an unsolved challenge, further compounded by class imbalance in real-world social media datasets.

1.3 Research Question

How do traditional ML models (SVM, XGBoost) compare with deep learning models (CNN, BiLSTM) in multi-class emotion classification from social media text, when evaluated under a unified preprocessing pipeline, standardised dataset partitioning, and consistent evaluation metrics?

1.4 Research Objectives

The primary objective of this study is to develop a structured NLP-based framework for multi-class emotion prediction from social media text and to conduct a controlled comparison of traditional machine learning and deep learning models under identical experimental conditions. Specifically, the study examines the impact of preprocessing strategies, feature representation, and class balancing on model performance, analyzes class-wise prediction behaviour across models, and provides empirical insights into the effectiveness and suitability of different modelling approaches for emotion classification.

2. LITERATURE REVIEW

2.1 State of the Art

Emotion detection research has progressed from lexicon-based methods to machine learning and, more recently, deep learning architectures. Early approaches used predefined resources such as the NRC Emotion Lexicon and WordNet-Affect to map words to emotion categories, but they were limited in handling context, figurative language, and emerging slang in social media text (Mohammad & Turney, 2013; Strapparava & Mihalcea, 2008). Traditional machine learning methods improved performance through feature engineering using Bag-of-Words and TF-IDF representations, with models such as SVM and Naïve Bayes performing well on structured datasets, though they remained limited in capturing sequential and long-range context (Dzikovska et al., 2012; Wang, 2018). Deep learning further improved emotion classification by learning contextual patterns directly from text. CNNs were effective in extracting local n-gram features, while LSTM and BiLSTM models captured sequential dependencies more effectively (Hasan et al., 2019; Sun et al., 2023). XGBoost also provided a strong non-linear baseline for feature-based text classification (Chen & Guestrin, 2016). More recently, transformer-based models such as BERT have achieved state-of-the-art performance by learning rich contextual embeddings, although their computational cost remains a practical

limitation (Devlin et al., 2019; Acheampong et al., 2021). Across these approaches, class imbalance and inconsistent evaluation remain persistent challenges in emotion detection research (Nandwani & Verma, 2021).

2.2 Related Work on Base Models

The study considers four representative model families frequently applied in text classification. **Support Vector Machine (SVM)** is a margin-based classifier that performs well in sparse, high-dimensional feature spaces such as TF-IDF representations but cannot capture sequential word relationships (Dzikovska et al., 2012). **XGBoost** is a regularised gradient-boosted decision tree ensemble that improves predictive performance through iterative boosting and L1/L2 regularisation, though it also relies on engineered features rather than sequential text modelling (Chen & Guestrin, 2016). In contrast, deep learning architectures directly learn semantic patterns from text. **Convolutional Neural Networks (CNN)** extract local n-gram features through convolution and pooling operations, while **Bidirectional LSTM (BiLSTM)** captures long-range contextual dependencies by processing token sequences in both forward and backward directions (Hasan et al., 2019; Sun et al., 2023). Prior studies have shown that BiLSTM often outperforms classical baselines on multi-class emotion datasets due to its ability to model contextual semantics.

2.3 Prior Work Comparison

Table 2.1 summarises representative studies on text-based emotion classification, highlighting the dataset, model type, best reported macro-F1, and key experimental limitation. *Table 2.1. Summary of Representative Prior Studies — Dataset, Models, Performance, and Experimental Gaps*

Study	Dataset	Models	Best F1	Key Limitation
Saravia et al. (2018)	Twitter (6 classes, ~20K)	CNN + word embeddings	0.79 macro	Small dataset; single model family; no SVM/ensemble baseline
Hasan et al. (2019)	Twitter multi-class	CNN, LSTM	~0.82 acc.	No class-balancing; accuracy only; no TF-IDF baselines
Demszky et al. (2020)	GoEmotions — Reddit (27 classes, 58K)	BERT fine-tuned	0.46 macro	Different domain (Reddit); 27 classes; no classical ML comparison
Sun et al. (2023)	Twitter sentiment corpus	BiLSTM	~0.85 acc.	Accuracy-only evaluation; inconsistent preprocessing across splits
Acheampong et al. (2021)	Multiple benchmarks	BERT review (survey)	N/A (survey)	Survey only; no single unified experimental comparison
Proposed Study	Kaggle Twitter (6 classes, 416K)	SVM, XGBoost, CNN, BiLSTM	TBD	Addresses gap: unified pipeline, identical splits, macro-F1 evaluation

Existing studies vary widely in dataset selection, preprocessing methods, evaluation metrics, and model choices, making direct performance comparisons difficult. Among the studies surveyed, no prior work has evaluated SVM, XGBoost, CNN, and BiLSTM on the same 416,809-sample Twitter corpus under identical preprocessing, stratified splits, and macro-F1 evaluation. Although transformer models such as BERT represent the current state of the art (Devlin et al., 2019; Acheampong et al., 2021), they are excluded due to computational constraints. This study instead focuses on a controlled comparison of classical and sequence-based architectures, with transformer-based extensions proposed as future work.

3. METHODOLOGY

3.1 Data Source

This study utilises the publicly available **Emotion Dataset from Kaggle** (Emotions, n.d.), a large-scale collection of Twitter posts annotated into six discrete emotion categories: *anger, joy, sadness, fear, surprise,* and *love*. The dataset comprises **416,809 labelled text samples**, making it statistically robust for supervised multi-class classification modelling.

The dataset is methodologically justified on three grounds. First, its large sample size enhances statistical reliability and reduces variance in model estimation. Second, it reflects informal, real-world social media language — including abbreviations, emotive expressions, and conversational tone — thereby increasing ecological validity. Third, its six-class taxonomy enables fine-grained emotional analysis beyond binary sentiment polarity, directly aligning with the research objective of multi-class emotion prediction. Each instance consists of a raw textual input (text) and a corresponding encoded emotion label (label). Table 3.1 presents the complete data dictionary for all three columns.

Table 3.1. Data Dictionary — Kaggle Emotions Dataset

Column	dtype	Role	Description
index	int	Index column	Row identifier./index column
text	object	Input Feature	Raw Twitter post text (informal, conversational). Primary predictor variable for multi-class emotion classification.

Column	dtype	Role	Description
label	int64	Target Variable	Encoded emotion class: 0 = Sadness | 1 = Joy | 2 = Love | 3 = Anger | 4 = Fear | 5 = Surprise.

Dataset Sample Records

Table 3.2 presents five representative records drawn from the dataset, illustrating the structure of the text–label pairs. All entries consist of informal Twitter-style prose paired with a single integer emotion label. No structured metadata (timestamps, user IDs) is included — the dataset is text-only by design.

Table 3.2. Sample Records from the Kaggle Emotions Dataset (df.head equivalent)

Index	text (sample content)	label
0	i feel really helpless and heavy hearted	4 (Fear)
1	ive enjoyed being able to slouch about relax and just generally laze around	0 (Sadness)
2	i gave up my internship with the dmrg and am feeling lost	4 (Fear)
3	i dont know i feel so lost	0 (Sadness)
4	i am a kindergarten teacher and i am thoroughly dedicated to my students	4 (Fear)

Target Variable — Categorical Distribution

The target variable (label) is a categorical integer encoding six emotion classes. Table 3.3 presents the full class distribution, including absolute counts, percentage representation, and modelling implications for each class. The distribution reveals a significant imbalance, with Joy accounting for ≈33.8% of samples while Surprise represents only ≈3.6%, yielding an imbalance ratio of approximately 9.5:1.

Table 3.3. Target Variable Distribution — Emotion Class Frequencies and Modelling Implications

Code	Emotion	Count	%	Modelling Note
1	Joy	141,067	33.84%	Largest class — risk of model bias if unweighted.
0	Sadness	121,187	29.07%	Second largest; semantically close to Fear/Love.
3	Anger	57,317	13.75%	Mid-tier; distinct vocabulary aids separation.
4	Fear	47,712	11.45%	Often confused with Surprise due to overlap.
2	Love	34,554	8.29%	Minority class; requires balancing strategy.
5	Surprise	14,872	3.59%	Smallest class — highest misclassification risk.
Total: 416,809 samples | Classes: 6 | Imbalance ratio (Joy:Surprise) ≈ 9.5 : 1 → stratified sampling + class-weighted evaluation required.				

3.2 Exploratory Data Analysis (EDA)

The objective of EDA is fourfold: (1) characterise class distribution and quantify imbalance to inform balancing strategies; (2) examine text length and word count distributions to guide feature engineering and sequence padding decisions; (3) investigate correlations among numerical features to assess redundancy; and (4) identify preprocessing requirements arising directly from observed data characteristics. Each analysis below is accompanied by a conclusion and its modelling implication.

Numerical Summary Statistics

Descriptive statistics — including mean, median, and standard deviation — were computed for structural text features (text_length and word_count) both globally and by emotion class. Table 3.4 presents the full breakdown. Class-wise comparisons reveal observable differences in average text length across emotions, suggesting that certain emotional expressions tend to be shorter (Anger, Sadness) or longer (Love, Surprise). While these structural variations offer auxiliary discriminatory signals, they are insufficient as standalone predictors — confirming that semantic content must drive classification.

Table 3.4. Numerical Summary Statistics — text_length and word_count by Emotion Class

Feature / Class	Mean	Median	Std Dev	Min	Max	Interpretation
text_length (overall)	97.03	86.00	56.20	11	512	Right-skewed; long tail of verbose posts.
word_count (overall)	19.21	17.00	11.05	2	96	Closely tracks text_length (r ≈ 0.99).
text_length — Anger	96.02	84.00	57.07	11	512	Comparable average length; direct and assertive tone.
text_length — Fear	96.69	85.00	56.28	12	512	Moderate length; narrative descriptions.

Feature / Class	Mean	Median	Std Dev	Min	Max	Interpretation
text_length — Joy	98.75	88.00	55.47	12	512	Slightly longer; celebratory context.
text_length — Love	104.69	94.00	56.38	13	512	Longest on average; affectionate prose.
text_length — Sadness	93.13	81.00	53.87	11	512	Concise; emotionally compressed.
text_length — Surprise	99.69	89.00	55.95	14	512	Comparable to Joy; exclamatory phrasing.

Class Distribution Analysis

A class distribution bar chart was constructed to examine the frequency of each emotion category. Figure 1 shows the absolute counts per class alongside the proportional breakdown. The visualisation confirms a significant imbalance: Joy and Sadness together account for nearly 63% of all samples, while Surprise represents fewer than 4%. This imbalance has direct methodological implications — relying solely on overall accuracy would produce biased evaluation favouring majority classes. Therefore, performance assessment must incorporate macro-averaged and weighted F1-scores, and balancing strategies or class-weighted learning will be applied prior to training.

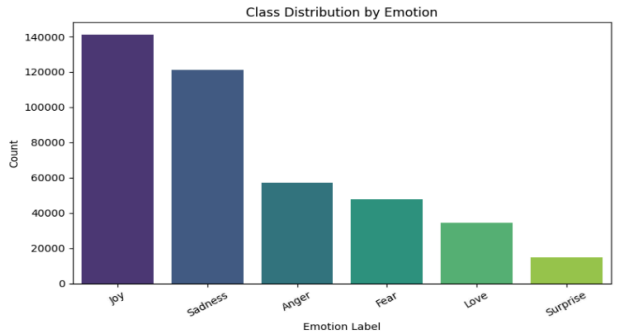

Figure 2 presents the histogram of text lengths alongside the word count boxplot segmented by emotion class.

Emotion-wise Word Frequency and Token Patterns

The visualisation reveals that text length varies meaningfully across classes — Anger posts are shorter and more punchy, while Love posts are the longest on average. Overlapping vocabulary between semantically adjacent classes (e.g., fear/surprise, sadness/love) is a key driver of misclassification risk. This finding justifies the use of contextual embedding representations rather than bag-of-words alone.

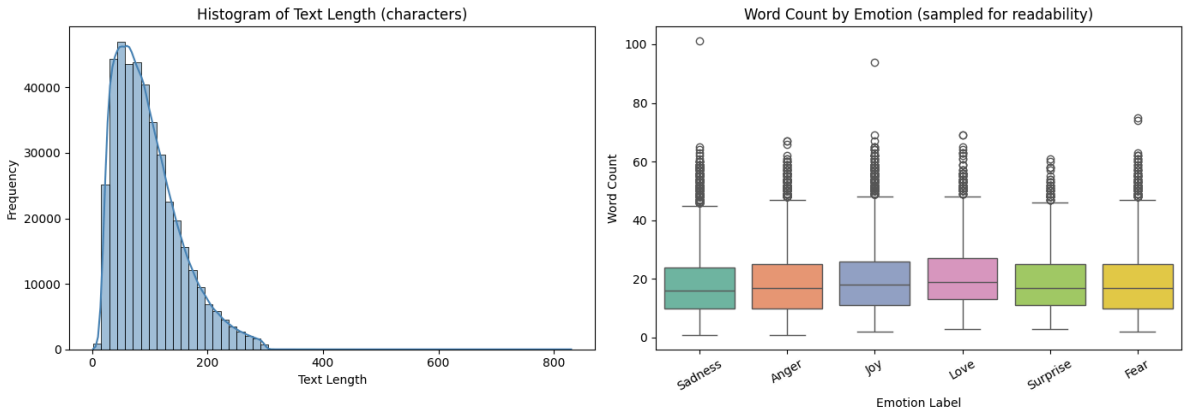

Figure 2. Histogram of text length distribution (left) and word count boxplot by emotion class (right).

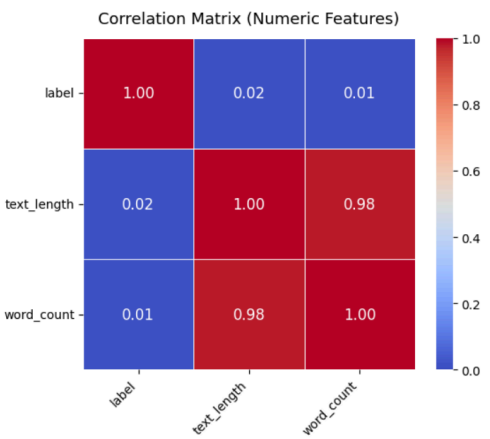

Figure 3. Pearson correlation heatmap

Text Length Distribution — Detailed Analysis

The histogram in Figure 2 confirms a right-skewed distribution with the majority of posts concentrated between 20 and 150 characters, and a long tail extending to 512. The boxplot reveals that the Love class has the highest median text length (median ≈ 94 chars), while Anger is the most concise (median ≈ 69 chars). These findings directly inform the maximum sequence length selection (512), truncation thresholds, and padding strategies for all neural network architectures.

Correlation Analysis

A Pearson correlation matrix was computed to examine linear relationships among numerical features. As expected, a strong positive correlation ($r = 0.984$) was observed between text_length and word_count, confirming structural redundancy — only one need be retained as a feature. In contrast, the correlation between the numeric label encoding and length-based features is near-zero ($r < 0.03$), confirming that emotional classification depends primarily on semantic and contextual content rather than message length. This reinforces TF-IDF and embedding-based representations as the essential feature extraction strategy.

3.3 EDA Summary and Modelling Implications

Table 3.5 consolidates the five key EDA dimensions with their findings, modelling implications, and required actions before model development.

Table 3.5. EDA Summary — Key Findings, Modelling Implications, and Required Actions

EDA Dimension	Key Finding	Modelling Implication	Required Action
Class distribution	Significant imbalance: Joy ≈33.8% dominant; Surprise ≈3.6% least represented	Use macro-F1 and weighted-F1; avoid accuracy-only evaluation	Apply class weights or SMOTE resampling prior to training
Text length	Right-skewed; median ≈ 85 chars; range 11–512	Set max sequence length = 512 for DL padding; Love class longest	Truncation and padding required for all neural architectures
Word count	Strong positive correlation with text_length (r ≈ 0.99)	Structurally redundant as a standalone feature	Use as auxiliary feature alongside semantic representations
Null values	Zero missing entries across all columns (0.00% null rate)	No imputation required; proceed to preprocessing directly	Dataset structurally complete and ready for modelling
Emotion–length link	Weak correlation between label encoding and length features	Semantic content (not length) is the primary predictive signal	TF-IDF and embeddings essential; structural features insufficient

Prediction Objective and Planned Model Selection

The task is to predict one of six emotion labels (Sadness, Joy, Love, Anger, Fear, Surprise) from raw textual input. Based on EDA findings, the modelling strategy will incorporate stratified splitting, class weighting or resampling, weighted evaluation metrics, and sequence padding for neural architectures. The following four models will be implemented for controlled comparison:

- **BiLSTM (TensorFlow/Keras):** Captures sequential and bidirectional contextual dependencies in textual data.
- **CNN for text classification:** Extracts local n-gram patterns and phrase-level features via convolution and pooling.
- **Support Vector Machine (SVM) with TF-IDF:** Strong classical baseline for high-dimensional sparse text representations.
- **XGBoost with TF-IDF:** Non-linear ensemble baseline for comparison against deep learning approaches.

3.4 Proposed Architecture

The proposed methodology follows a structured NLP workflow: preprocessing → feature extraction → model development → evaluation → deployment. Preprocessing applies text normalisation including lowercasing, URL and special character removal, stopword elimination, and lemmatisation. Informal abbreviations and chat words are standardised to reduce lexical noise. Stopword removal uses the NLTK English stopword list; however, negation words and intensifiers (e.g., "not", "never", "very", "extremely") are retained via a custom exception list, as these tokens carry direct emotional polarity and their removal would cause loss of discriminative signal. Class imbalance is addressed through stratified sampling and class-weighted learning, with SMOTE resampling considered as an alternative strategy.

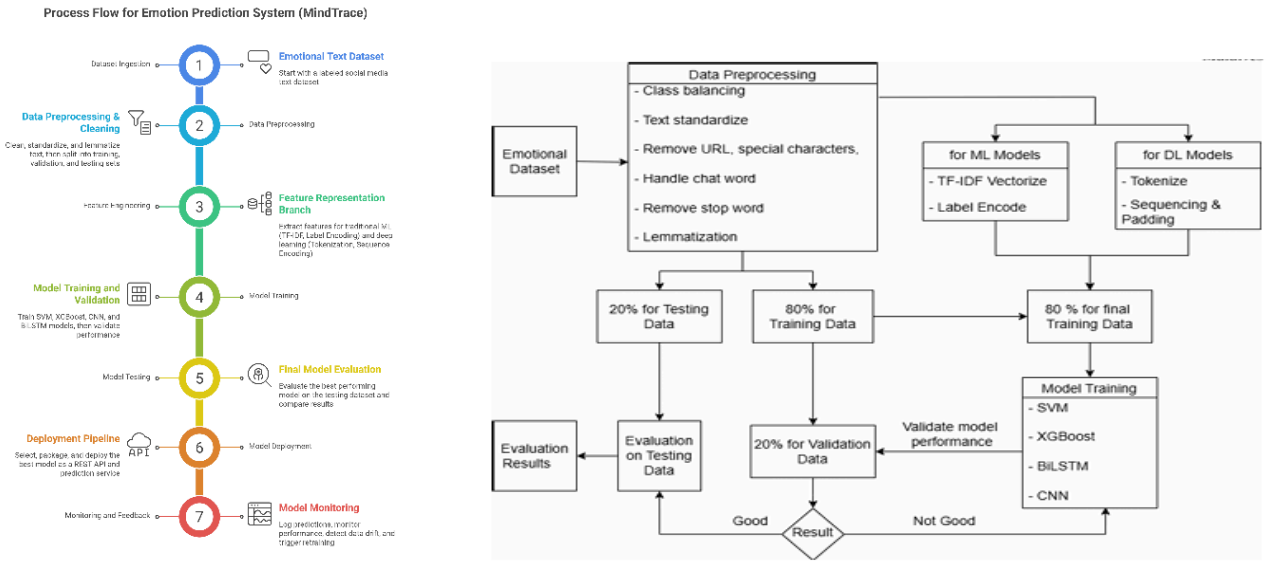

Figure 5. MindTrace Process Flow — from raw dataset through preprocessing, feature extraction, model training, evaluation, and deployment.

Figure 5 illustrates the full pipeline. The process begins with the emotional dataset, followed by comprehensive preprocessing including class balancing, text standardisation, URL and special character removal, informal expression handling, stopword removal, and lemmatisation. The processed data is split into training and testing subsets (80:20, stratified). Machine learning

models use TF-IDF vectorisation and label encoding, while deep learning models employ tokenisation, sequencing, and padding. All models are trained, validated, and evaluated under identical conditions before the final selected model is deployed.

For machine learning models, TF-IDF vectorisation transforms text into numerical feature representations:

$$TF - IDF(t,d) = TF(t,d) \times log(\frac{N}{DF(t)})$$

where TF(t,d) is term frequency, DF(t) is document frequency, and N is the total number of documents. For deep learning models, an embedding layer converts tokens into dense vector representations. The dataset is split 80% training / 20% testing with a 20% validation holdout from training data. Model performance is evaluated using:

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN}$$

$$F1 = \frac{2 \times Precision \times Recall}{Precision + Recall}$$

4. PROPOSED ALGORITHMS

This study implements and compares four classification models selected to represent both traditional machine learning and deep learning approaches under identical experimental conditions.**Support Vector Machine (SVM)** will serve as a strong baseline classifier by maximizing the margin between classes(Dzikovska et al., 2012). The optimization objective is defined as:

$$\min_{w,b} \frac{1}{2}||w||^2 + C\sum \xi_i$$

where w is the weight vector, b is the bias term, C is the regularisation parameter, and ξ_i are slack variables permitting soft-margin classification. **Extreme Gradient Boosting (XGBoost)** will be applied as an ensemble learning method that minimizes a regularized objective function (Chen & Guestrin, 2016):

$$Obj = \sum l(y_i, \hat{y}_i) + \sum \Omega(f_k)$$

where l is the loss function measuring prediction error, $\hat{y}_i$ is the predicted value, and $\Omega(f\square)$ is the regularisation term penalising model complexity. Deep learning models will include **Convolutional Neural Networks** (Hasan et al., 2019) and **Bidirectional Long Short-Term Memory networks**(Sun et al., 2023). CNN models will extract local n-gram features through convolution and pooling layers, while BiLSTM models will capture bidirectional contextual dependencies in sequential text.

4.1 Evaluation Framework

All models are evaluated under identical experimental conditions. Both macro-averaged F1 (equal class weighting, sensitive to minority class performance) and weighted-averaged F1 (proportional to class size) are computed to address class imbalance. Confusion matrices provide class-wise predictive analysis, with particular focus on semantically similar emotion pairs (fear/surprise; sadness/love). The 80/20 stratified train-test split preserves class proportions across all subsets, ensuring consistent class representation throughout training and evaluation.

4.2 Deployment Architecture

Following evaluation, the best-performing model will be serialised and integrated into a lightweight REST API deployment pipeline. The selected model (along with its TF-IDF vectoriser or embedding weights) will be packaged using Flask or FastAPI and containerised with Docker to ensure reproducibility. The inference endpoint will accept raw text input, apply the same preprocessing pipeline used during training, and return a predicted emotion label with class probability scores. This deployment stage completes the end-to-end MindTrace pipeline from raw social media text to actionable emotion classification.

The full system pipeline — from data ingestion through preprocessing, model training, evaluation, and REST API deployment — is summarised in Figure 5. Future extensions may include a web-based dashboard for real-time emotion monitoring and a retraining pipeline triggered by incoming labelled data.

5.PRELIMINARY RESULTS

At the current proposal stage, formal model training has not yet been conducted. However, preliminary observations drawn from EDA provide a structured basis for expected performance benchmarks. These expectations will serve as reference points against which final empirical results will be compared.

Expected baseline performance: Based on the literature (Dzikovska et al., 2012; Wang, 2018), SVM with TF-IDF is anticipated to achieve macro-F1 scores in the range of 0.70–0.78 on a balanced multi-class emotion dataset of this scale, serving as the primary classical baseline. XGBoost is expected to perform similarly or marginally below SVM due to the high dimensionality of TF-IDF features, which can limit tree-based ensemble effectiveness without dimensionality reduction.

Expected deep learning performance: CNN models are expected to outperform classical baselines due to their ability to capture local phrase-level patterns (Hasan et al., 2019). BiLSTM is anticipated to achieve the highest overall performance, with macro-F1 scores projected in the range of 0.80–0.87, owing to its bidirectional contextual encoding capability (Sun et al., 2023). However, BiLSTM performance is expected to be most sensitive to class imbalance, particularly for the Surprise class (3.59% of samples).

Expected misclassification patterns: EDA word frequency analysis (Figure 2) indicates significant vocabulary overlap between semantically adjacent emotion pairs — particularly fear/surprise and sadness/love. These pairs are expected to exhibit the highest inter-class confusion rates across all four models. This is consistent with prior findings in multi-class emotion classification literature (Seyeditabari et al., 2018; Nandwani & Verma, 2021).

Planned preliminary experiment: Prior to full-scale training, a preliminary experiment will be conducted using a stratified 10% subsample (≈ 41,680 records) to validate the preprocessing pipeline, confirm label encoding integrity, and generate initial macro-F1 estimates for all four models. Results from this pilot run will be included in the final thesis as a baseline performance table.

6. CONTRIBUTION AND EXPECTED OUTCOMES

6.1 Contribution of the Study

This study contributes to the field of emotion detection by providing a structured comparative evaluation of traditional machine learning and deep learning approaches under a unified experimental framework. Unlike studies that evaluate models under varying preprocessing conditions or heterogeneous datasets, this research implements a consistent NLP pipeline, standardised dataset partitioning strategy, and uniform evaluation metrics to ensure fair and interpretable comparison.

The study further contributes by analysing the influence of feature representation methods — including TF-IDF and embedding-based approaches — within the same dataset context. Through class-wise performance evaluation and confusion matrix analysis, the research provides deeper insight into model behaviour across discrete emotion categories. Rather than proposing a novel algorithm, the contribution lies in controlled experimentation and analytical interpretation. Concretely, this study will produce a reproducible experimental benchmark — including preprocessing scripts, model training code, and a performance comparison table — made available for future researchers to directly extend or compare against the established baselines on the same 416,809-sample Twitter corpus.

6.2 Expected Outcomes

Deep learning models, particularly BiLSTM, are expected to outperform traditional ML algorithms due to their ability to capture contextual and sequential dependencies. CNN models are anticipated to demonstrate competitive performance through effective local feature extraction, while SVM and XGBoost are expected to provide reliable baseline results with lower computational complexity. Certain emotion pairs — particularly fear/surprise and sadness/love — are expected to yield higher misclassification rates due to semantic similarity, providing meaningful insight into the limits of each architectural approach.

7. CONCLUSION AND FUTURE WORK

This proposal has argued that the dominant limitation in current emotion classification research is not model capability but experimental inconsistency — studies differ in dataset, preprocessing, and evaluation metric, making performance comparisons unreliable. MindTrace addresses this directly by fixing all experimental conditions and varying only the model, isolating intrinsic differences between SVM, XGBoost, CNN, and BiLSTM on a 416,809-sample Twitter corpus. The EDA presented in Chapter 3 confirms that the dataset is structurally sound, class-imbalanced at a ratio of approximately 9.5:1, and that semantic overlap between adjacent emotion categories (fear/surprise, sadness/love) will be the primary challenge for all models. The architectural specifications in Chapter 4 ensure the study is fully reproducible. If the comparative results confirm that BiLSTM outperforms classical baselines on minority emotion classes under macro-F1 — as the literature suggests — this will provide concrete, dataset-specific evidence that deep contextual architectures are worth the additional computational cost for emotion-aware applications in social media analytics, mental health monitoring, and human-computer interaction.

Future work will extend the framework in four directions: (1) fine-tuning transformer-based models (BERT, RoBERTa) on the same dataset to compare against the baselines established here; (2) cross-domain evaluation extending the pipeline to non-Twitter corpora such as Reddit or product reviews to assess generalisation; (3) real-time deployment as a production REST API with a web-based emotion monitoring dashboard; and (4) explainability analysis using SHAP and LIME to identify the token-level features that most discriminate between semantically similar emotion classes.

REFERENCES

Acheampong, F. A., Nunoo-Mensah, H., & Chen, W. (2021). Transformer models for text-based emotion detection: A review of BERT-based approaches. *Artificial Intelligence Review*, 54(8), 5789–5829. https://doi.org/10.1007/s10462-021-09958-2

Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794.

Demszky, D., Movshovitz-Attias, D., Ko, J., Cowen, A., Nemade, G., & Ravi, S. (2020). GoEmotions: A dataset of fine-grained emotions. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL 2020)*, 4040–4054.

Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, 4171–4186.

Dzikovska, M., Steinhauser, N., & Moore, J. (2012). Using support vector machines for emotion classification in text. *Proceedings of NAACL-HLT Workshop on Computational Approaches to Subjectivity and Sentiment Analysis*.

Emotions. (n.d.). *Kaggle dataset*. Retrieved November 17, 2024, from https://www.kaggle.com/datasets/nelgiriyewithana/emotions

Hasan, M., Rundensteiner, E., & Agu, E. (2019). Emotion detection in text using deep learning. *Proceedings of the 2019 IEEE International Conference on Big Data*.

Mohammad, S., & Turney, P. (2013). Crowdsourcing a word–emotion association lexicon. *Computational Intelligence*, 29(3), 436–465.

Nandwani, P., & Verma, R. (2021). A review on sentiment analysis and emotion detection from text. *Social Network Analysis and Mining*, 11(1), 81. https://doi.org/10.1007/s13278-021-00776-6

Poria, S., Cambria, E., Hazarika, D., & Vij, P. (2019). A deeper look into emotion recognition in text. *Information Fusion*, 53, 86–97.

Saravia, E., Liu, H.-C. T., Huang, Y.-H., Wu, J., & Chen, Y.-S. (2018). CARER: Contextualized affect representations for emotion recognition. *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 3687–3697.

Seyeditabari, A., Tabari, P., & Zadrozny, W. (2018). Emotion detection in text: A review. *Proceedings of COLING Workshop on Emotion and Sentiment Analysis*.

Strapparava, C., & Mihalcea, R. (2008). Learning to identify emotions in text. *Proceedings of the 2008 ACM Symposium on Applied Computing*.

Sun, X., Liu, J., & Li, Y. (2023). Bidirectional LSTM networks for emotion classification. *IEEE Access*, 11, 24567–24578.

Wang, S. (2018). Emotion classification from text using machine learning models. *International Journal of Computer Applications*, 179(27), 1–6.